


(a)

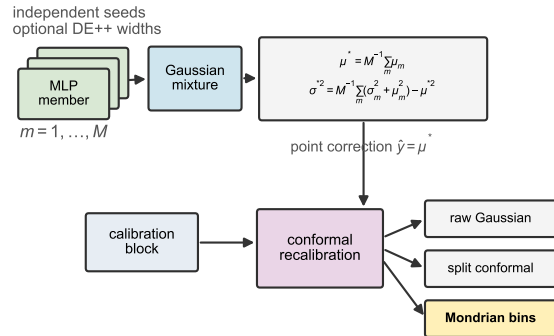
One Gaussian-head MLP member



Each hidden block is Linear -> LayerNorm -> ReLU -> Dropout ($p = 0.1$); dropout is active in training only, so there is no MC-dropout at inference. The second head output is clamped to $[-10, 10]$. For Student-t loss it is interpreted as log scale; otherwise it is log variance.

(b)

Ensemble mixture and conformal calibration



All interval variants share the same ensemble mean. Mondrian conformal uses per-bin residual quantiles from predicted-mean deciles, so no cloud information is needed at inference.